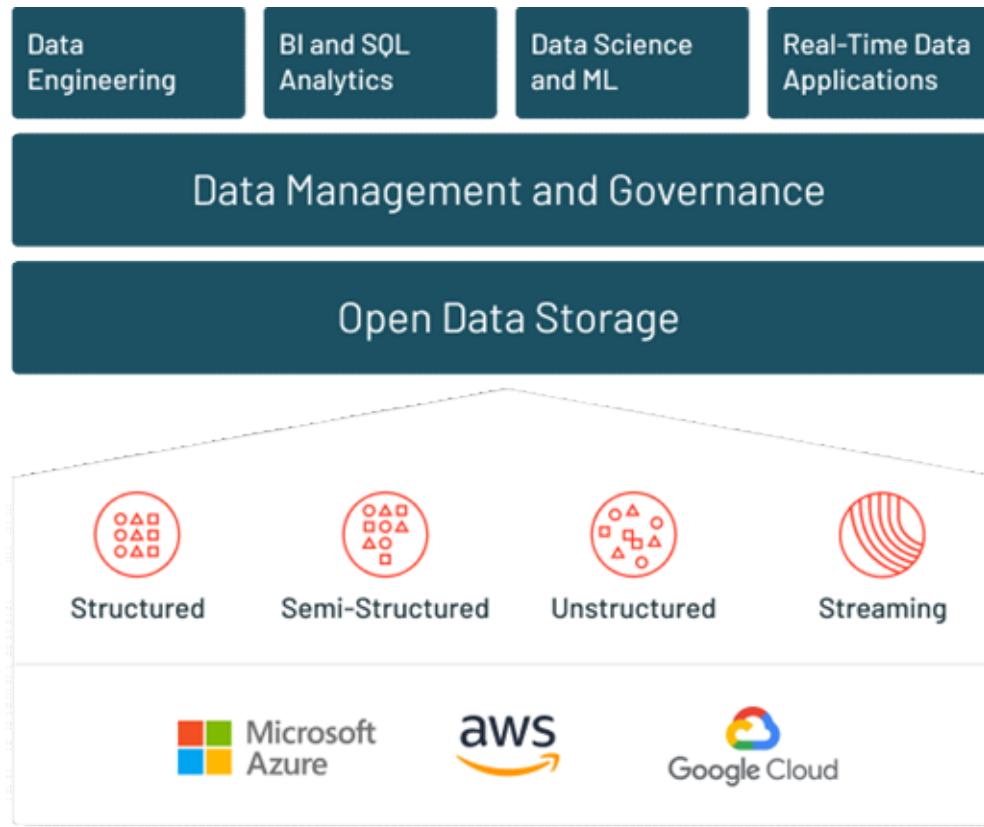



DATABRICKS on AWS

What is Databricks ?

- **Databricks** is a cloud-based platform for managing and analyzing large datasets using Apache Spark.



What is Databricks ?



Databricks features

- **Unified Workspace**

- Databricks provides a single platform for data scientists, engineers, and business analysts to work together and collaborate on data projects. This can help to improve communication and collaboration within teams and make it easier to develop and deploy data-driven applications.

- **Scalability and Flexibility**

- Databricks is designed to be highly scalable so that they can easily handle large amounts of data. It can also be flexibly configured to support different workloads, such as batch processing, real-time streaming, or machine learning. This makes it a good choice for organizations that need to process data at different scales, or that have complex data pipelines.

- **Integrated Tools and Services**

- Databricks comes with a range of tools and services for working with big data, including support for various data formats, integration with popular data science libraries and frameworks, and tools for data preparation and analysis. This makes it easier to build and deploy data-driven applications, without having to worry about setting up and managing complex infrastructure.

- **Security and Compliance**

- Databricks offers a range of features to help ensure that data is handled securely and in compliance with relevant regulations. This includes support for encryption, role-based access control, and auditing, as well as integration with popular security and compliance tools.

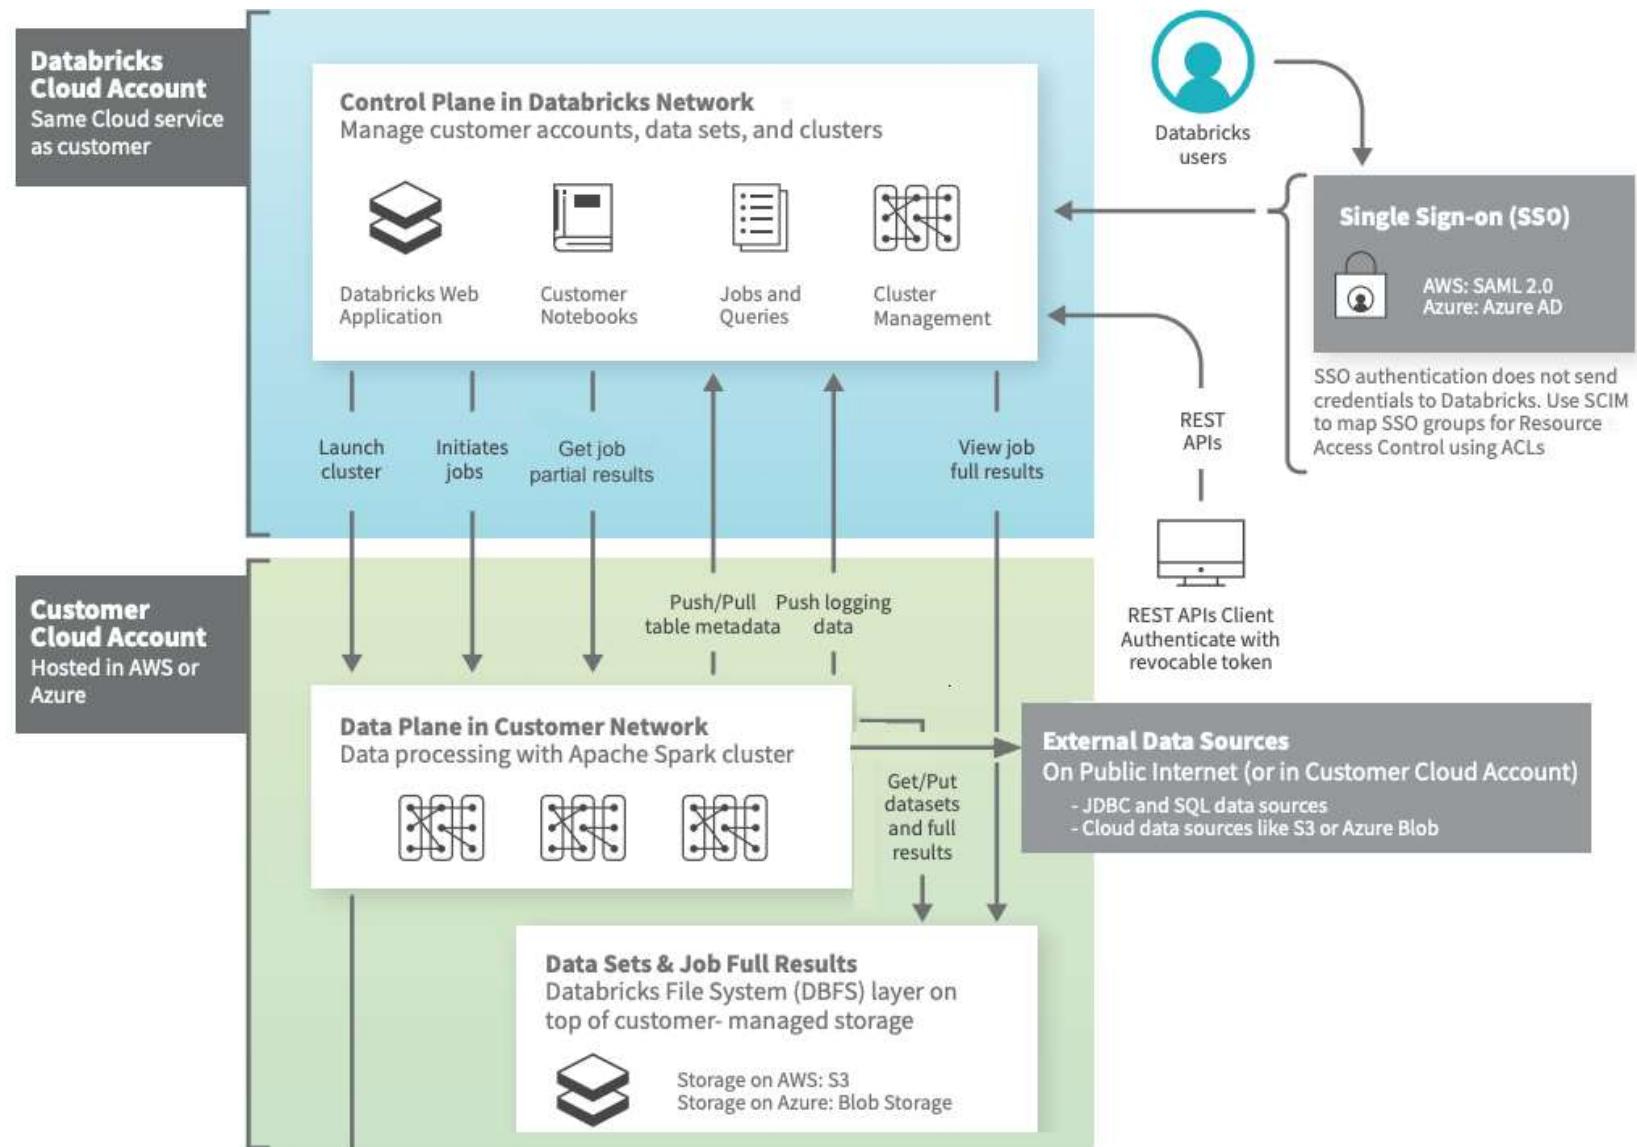
Databricks Use Cases

- **Data Warehousing**
 - Databricks can be used to store and manage large amounts of data from multiple sources, and provide fast and efficient access to the data for analysis and other purposes.
- **Data Preparation**
 - Databricks provides tools and services for cleaning, transforming, and enriching data, making it easier to prepare data for analysis or other uses.
- **Data Analysis**
 - Databricks offers a range of tools and services for exploring, visualizing, and analyzing data, so that users can gain insights and identify trends in the data.
- **Machine Learning**
 - Databricks supports popular machine learning libraries and frameworks, such as TensorFlow, Keras, and PyTorch, making it easier to build and train ML models on large datasets.
- **Real-time Data Processing**
 - Databricks can be used to process streaming data in real-time so that users can take action based on the data as it arrives.

Databricks Architecture

- Databricks is structured to enable secure cross-functional team collaboration while keeping a significant amount of backend services managed by Databricks so you can stay focused on your data science, data analytics, and data engineering tasks.
- Databricks operates out of a ***control plane*** and a ***data plane***.
 - **Control plane**
 - The control plane includes the backend services that Databricks manages in its own AWS account. Notebook commands and many other workspace configurations are stored in the control plane and encrypted at rest.
 - **Data plane**
 - For most Databricks computation, the compute resources are in your AWS account in what is called the Classic data plane. This is the type of data plane Databricks uses for notebooks, jobs, and for pro and classic Databricks SQL warehouses.
- Your data lake is stored at rest in your own AWS account.
- Job results reside in storage in your account.

Databricks Architecture



Databricks Workspace

- A Databricks workspace is a software-as-a-service (SaaS) environment for accessing all Databricks assets.
- The workspace organizes objects (for example, notebooks, libraries, and experiments) into folders and provides access to data and computational resources, such as clusters and jobs.

AWS CLI

- The AWS Command Line Interface (CLI) is a unified tool to manage your AWS services.
- With just one tool to download and configure, you can control multiple AWS services from the command line and automate them through scripts.
- The AWS CLI v2 offers several new features including improved installers, new configuration options such as AWS Single Sign-On (SSO), and various interactive features.

AWS CLI

- The AWS Command Line Interface (AWS CLI) is an open source tool that enables you to interact with AWS services using commands in your command-line shell.
- With minimal configuration, the AWS CLI enables you to start running commands that implement functionality equivalent to that provided by the browser-based AWS Management Console from the command prompt in your terminal program:
 - **Linux shells** – Use common shell programs such as bash, zsh, and tcsh to run commands in Linux or macOS.
 - **Windows command line** – On Windows, run commands at the Windows command prompt or in PowerShell.
 - **Remotely** – Run commands on Amazon EC2 instances through a remote terminal program such as PuTTY or SSH, or with AWS Systems Manager.

Databricks Cluster

- A Databricks cluster is a set of computation resources and configurations on which you run data engineering, data science, and data analytics workloads, such as production ETL pipelines, streaming analytics, ad-hoc analytics, and machine learning.
- You run these workloads as a set of commands in a **notebook** or as an **automated job**.
 - Databricks makes a distinction between all-purpose clusters and job clusters.
 - All-purpose clusters are used to analyze data collaboratively using interactive notebooks.
 - Job clusters are used to run fast and robust automated jobs.

Databricks cluster types

There are mainly two types of clusters in Databricks:

- **All-Purpose Compute**

- These are mainly used to analyze data interactively using Databricks notebooks.
- We can create these clusters using the Databricks UI, CLI, or REST API commands and also, can manually stop and restart these clusters.
- Multiple users can share these clusters to do collaborative interactive analysis.

- **Job Compute**

- Databricks job scheduler creates these clusters when we run a job on a new job cluster.
- These are mainly used for running fast and robust automated tasks. They are created when we run a job on your new Job Cluster and terminate the Cluster once the job ends.
- These clusters cannot be restarted.

Databricks cluster policies

Personal Compute

- The Personal Compute policy is available to all users in your workspace by default. This policy allows users to easily create a single-node compute resource with minimal configuration options.

Power User Compute

- The Power User Compute allows users to create larger multi-node resources. This is intended for single-user workloads that require more compute resources than Personal Compute allows.
- Key details for this policy include:
 - All-purpose compute only (no jobs)
 - Uses the standard access mode to allow multiple users to use this resource
 - By default, it can only be used by workspace admins

Shared Compute

- The Shared Compute policy allows users to create larger multi-node resources intended for multiple users to share.
- Key details for this policy include:
 - All-purpose compute only (no jobs)
 - Uses the standard access mode to allow multiple users to use this resource
 - Defaults to latest LTS Databricks Runtime version
 - By default, it can only be used by workspace admins

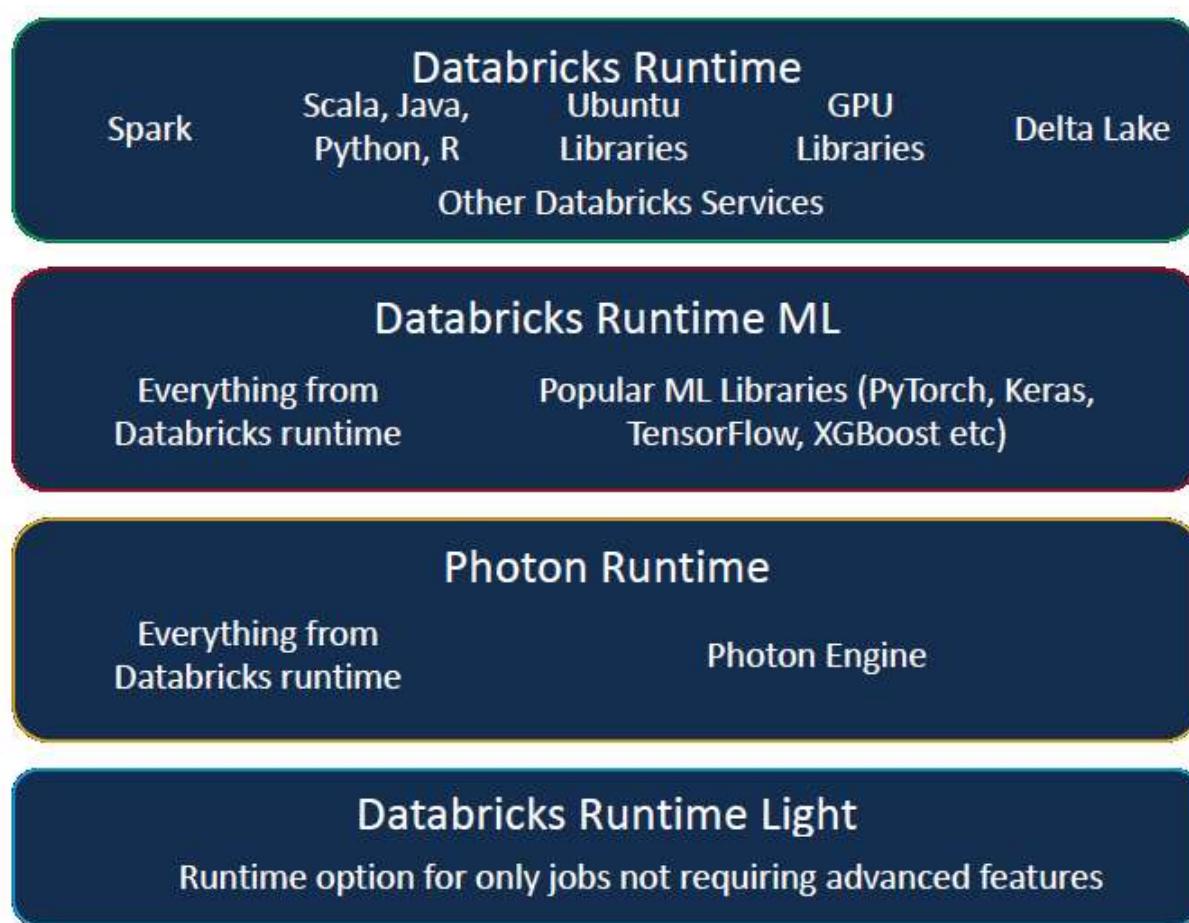
Databricks cluster access modes

Cluster access mode is a security feature that determines who can use a cluster and what data they can access via the cluster.

Access Mode	Visible to user	UC Support	Supported Languages	Notes
Single user	Always	Yes	Python, SQL, Scala, R	Can be assigned to and used by a single user. See single user limitations .
Shared	Always (Premium plan required)	Yes	Python (on Databricks Runtime 11.1 and above), SQL	Can be used by multiple users with data isolation among users. See shared limitations .
No Isolation Shared	Admins can hide this cluster type by enforcing user isolation in the admin settings page.	No	Python, SQL, Scala, R	There is a related account-level setting for No Isolation Shared clusters.
Custom	Hidden (For all new clusters)	No	Python, SQL, Scala, R	This option is shown only if you have existing clusters without a specified access mode.

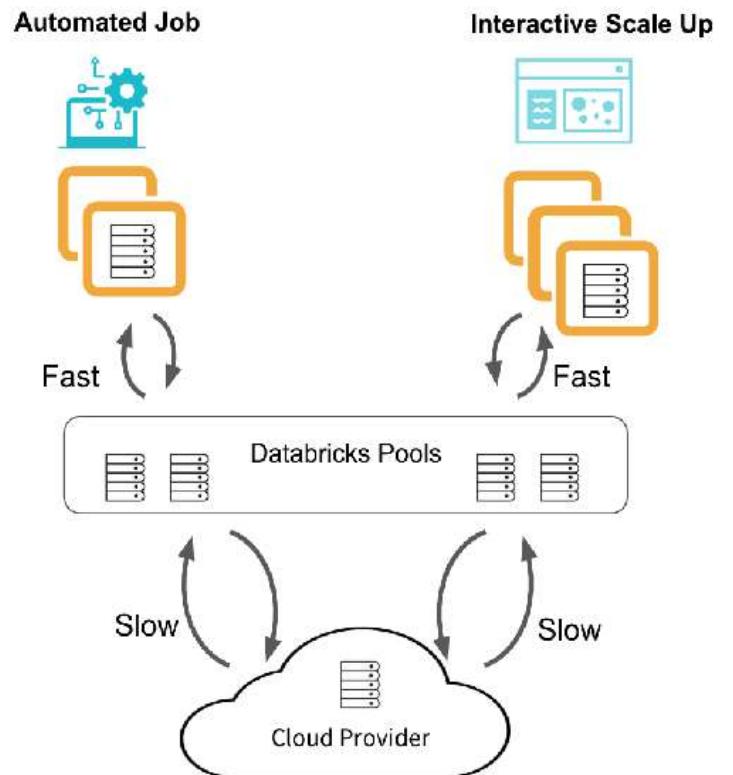
Databricks runtime options

Databricks Runtime is the set of software artifacts that run on the clusters of machines managed by Databricks.



Databricks cluster pools

- Azure Databricks pools are a set of idle, ready-to-use instances.
- When cluster nodes are created using the idle instances, cluster start and auto-scaling times are reduced.
- If the pool has no idle instances, the pool expands by allocating a new instance from the instance provider in order to accommodate the cluster's request.
- When a cluster releases an instance, it returns to the pool and is free for another cluster to use. Only clusters attached to a pool can use that pool's idle instances.



Create pools based on workloads

- If your driver node and worker nodes have different requirements, create a different pool for each.
- You can minimize instance acquisition time by creating a pool for each instance type and Azure Databricks runtime your organization commonly uses.
 - For example, if most data engineering clusters use instance type A, data science clusters use instance type B, and analytics clusters use instance type C, create a pool with each instance type.
- Configure pools to use on-demand instances for jobs with short execution times and strict execution time requirements. Use on-demand instances to prevent acquired instances from being lost to a higher bidder on the spot market.
- Configure pools to use spot instances for clusters that support interactive development or jobs that prioritize cost savings over reliability.

Using the pool for your cluster

Clusters are attached to the pool. You can specify a different pool for the driver node and worker nodes, or use the same pool for both.

Multi node Single node

Access mode ? Single user access ?

Single user | ▼ Kanakaraju Yarakaraju (kanakaraj... | ▼

Performance

Databricks runtime version ?

Runtime: 12.2 LTS (Scala 2.12, Spark 3.3.2) | ▼

Use Photon Acceleration ?

Worker type ?	Min workers	Max workers	
CTS Demo Pool	Standard_F4 ▼	2	8

Databricks Photon

Databricks Photon is a vectorized query engine designed to accelerate Spark workloads on the Databricks Lakehouse Platform. It is not a replacement for Spark, but rather an optimized execution engine that works within the Spark framework.

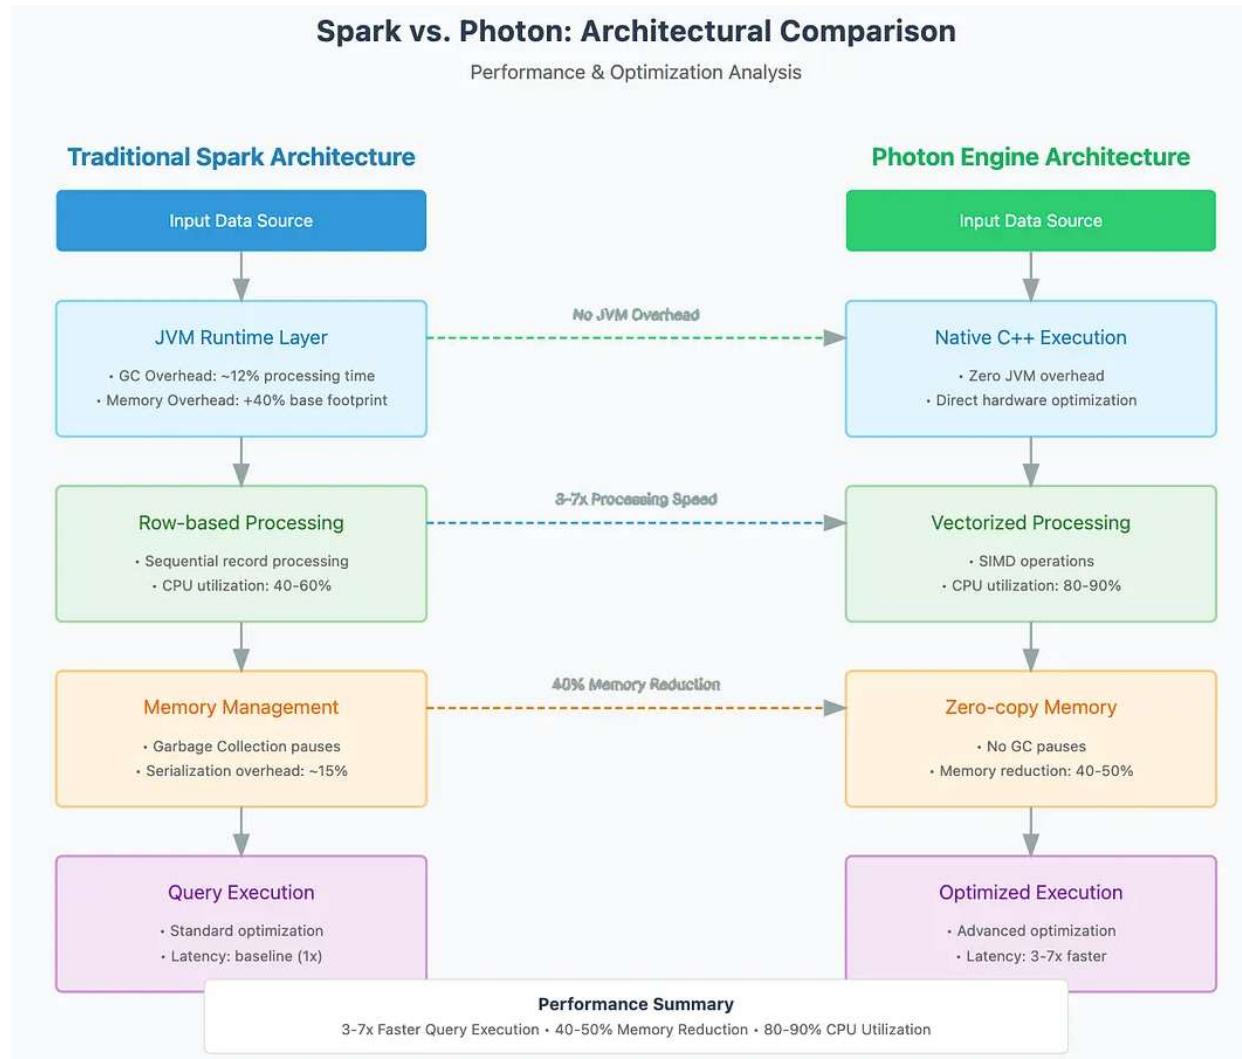
- **Execution Language:** Reimplemented in C++ for native execution.
- **Execution Model:** Vectorized execution, processing data in batches (columns) rather than row by row. This amortizes the cost of operations across multiple values.
- **Performance:** Designed for significant speedups, especially for SQL queries and DataFrame transformations, leveraging modern CPU architectures and techniques like AVX-512 instruction sets. It aims to reduce CPU cycles and improve memory bandwidth utilization.
- **Integration:** Works seamlessly with existing Spark APIs and SQL queries, utilizing Spark's Catalyst optimizer for query planning. Photon essentially replaces the underlying execution engine for compatible operations.
- **Cost:** While offering performance benefits, Photon typically incurs a higher Databricks Unit (DBU) cost compared to traditional Spark execution, requiring a cost-benefit analysis.

Databricks Photon

Feature	Apache Spark (Traditional)	Databricks Photon
Execution Engine	JVM-based (Scala/Java)	C++ native engine
Execution Model	Row-by-row processing	Vectorized (batch) processing
Performance	Good, but can be impacted by JVM overhead	Significantly faster, especially for analytical workloads
Compatibility	Fully compatible with Spark APIs	Works with Spark APIs and SQL, but replaces the execution engine for compatible operations
Cost	Standard DBU pricing	Higher DBU cost due to enhanced performance

In essence, Photon enhances Spark's capabilities by providing a highly optimized, native execution engine for data processing, leading to faster query execution and potentially lower overall compute costs if the speedup outweighs the increased DBU cost.

Databricks Photon



Databases & Tables

- Databricks Database is a collection of tables.
- Databricks Table is a collection of structured data.
 - We can cache, filter, and perform any operations supported by Spark DataFrames on Azure Databricks tables and query tables with Spark APIs and Spark SQL.
 - There are two types of tables.
 1. Global Table
 - A global table is available across all clusters.
 - Databricks registers global tables to either the Hive metastore or an external metastore.
 - These are persistent tables (i.e. not temporary tables)
 2. Local Table
 - A local table is not accessible from other clusters
 - Not registered in the Hive metastore. (i.e. not created in any database)
 - This is also known as a temporary view.

Volumes


- A volume is a Unity Catalog object that provides a logical layer for storing, governing, and accessing non-tabular data, such as images, audio, or PDFs, in any format. Volumes allow you to manage files with centralized governance.
- You can create two types of volumes:
 - **Managed volumes** - Databricks manages the underlying cloud storage location and lifecycle.
 - **External volumes** - You specify the cloud storage location and have control over the lifecycle.

```
CREATE VOLUME <catalog>.<schema>.<volume-name>
```

```
CREATE EXTERNAL VOLUME <catalog>.<schema>.<volume-name>
LOCATION '<cloud-uri>' ;
```

Databricks CLI

- The Databricks command-line interface (Databricks CLI) utility provides an easy-to-use interface to automate the Databricks platform from your terminal, command prompt, or automation scripts.



A screenshot of an Anaconda Prompt window. The title bar says "Anaconda Prompt". The command line shows the user running the Databricks CLI. The first command is "dbcli-venv) (base) E:\Databricks\Projects> databricks -v", which outputs "Version 0.17.7". The second command is "dbcli-venv) (base) E:\Databricks\Projects> databricks fs ls --profile ctsdemo", which lists several Databricks file system paths: "FileStore", "Volume", "Volumes", "databricks-datasets", "databricks-results", "jobs", "mnt", "tmp", "user", "volume", and "volumes".

```
(dbcli-venv) (base) E:\Databricks\Projects> databricks -v
Version 0.17.7

(dbcli-venv) (base) E:\Databricks\Projects> databricks fs ls --profile ctsdemo
FileStore
Volume
Volumes
databricks-datasets
databricks-results
jobs
mnt
tmp
user
volume
volumes
```

Databricks File System (DBFS)

- The Databricks File System (DBFS) is a distributed file system mounted into a Databricks workspace and available on Databricks clusters.
- DBFS is an abstraction on top of scalable object storage that maps file-system calls to native cloud storage API calls.

What can you do with DBFS?

- DBFS provides convenience by mapping cloud object storage URIs to relative paths.
- Allows you to interact with object storage using directory and file semantics instead of cloud-specific API commands.
- Allows you to mount cloud object storage locations so that you can map storage credentials to paths in the Databricks workspace.
- Simplifies the process of persisting files to object storage, allowing virtual machines and attached volume storage to be safely deleted on cluster termination.
- Provides a convenient location for storing init scripts, JARs, libraries, and configurations for cluster initialization.
- Provides a convenient location for checkpoint files created during model training with OSS deep learning libraries.



Introduction to Delta Lake

Data Warehouse, Data Lake & Lakehouse Concepts

Relational databases

- In 1970, EF Codd introduced the concept of looking at data as logical relations, independent of the physical data storage. This logical relation between data entities became known as a *database model* or *schema*. Codd's writings led to the birth of the ***relational database***.
- Relational databases and their underlying SQL language became the standard storage technology for enterprise applications throughout the 1980s and 1990s. One of the main reasons behind this popularity was that relational databases offered a concept called ***transactions***.
- A database transaction is a sequence of operations on a database that satisfies the ACID properties, ACID is an acronym which stands for four properties:
 - Atomicity
 - Consistency
 - Isolation
 - Durability


ACID properties

- **Atomicity** ensures that all changes made to the database are executed as a single operation.
 - For example, when I use my online banking to transfer money, the atomicity property will guarantee that the operation will only succeed when the money is deducted from one account and added to the other. The complete operation will either succeed or fail as a unit.
- **Consistency** property guarantees that database transitions from one consistent state of the transaction to another consistent state at the end of the transaction.
 - The transfer of the money would only happen if my savings account had sufficient funds. If not, the transaction would fail, and the balances would stay in their original, consistent state
- **Isolation** ensures that concurrent operations that are happening within the database are not affecting each other.
 - This property ensures that when multiple transactions are executed concurrently, their operations do not interfere with each other
- **Durability** refers to the persistence of committed transactions. It guarantees that once a transaction is completed successfully, it will result in a permanent state even in the event of a system failure.
 - Durability will ensure that updates made to both my savings and checking account are persistent and can survive a potential system failure.

Data warehouse

- Enterprise applications were using the RDBMS technology very effectively. Flagship products such as SAP and Salesforce would collect and maintain massive amounts of data.
- Enterprise applications would store the data in their own, proprietary formats, leading to the rise of ***data silos***. These data silos were owned and controlled by one department or business unit.
- Over time, organizations recognized the need to develop an **Enterprise view across these different data silos**, leading to the rise of ***data warehouses***.
- **Data warehouse** is a **central relational repository** of integrated, historical data from **multiple data sources** that presents a single integrated, historical view of the business with a **unified schema, covering all perspectives of the enterprise**.

Data warehouse




The 'Big Data' problem

- The fast rise of the internet, social media and the availability of multi-media devices such as smart phones disrupted the traditional data landscape, giving rise to the term **big data**.
- Big data is defined as data that arrives in ever higher **volumes**, with more **velocity**, and a greater **variety** of formats and has higher **veracity**. These are known as the four Vs of data.
- Data Warehouses have a hard time addressing these four Vs.


Data Lake

- A **Data Lake** is a *cost-effective central repository to store structured, semi-structured or unstructured data at any scale*, in the form of files and blobs.
- The initial data lakes and big data solutions were built with on-premises clusters, based upon the Apache **Hadoop** open-source set of frameworks. Hadoop was used to efficiently store and process large datasets ranging in GBs to PBs data.
- Starting in 2015, **cloud data lakes** such as **S3, ADLS, and GCS** started replacing HDFS. These cloud-based storage systems have superior SLAs, offer geo-replication and most importantly, offer extremely low cost with the option to utilize even lower-cost cold storage for archival purposes.




Data Lake & Data Warehouse

Data Lake



Data Warehouse



Benefits of Data Lake

- A data lake architecture enables the consolidation of an organization's data assets into one central location.
- They use open formats, such as Parquet and Avro. These formats enable easy interoperability, especially in our open-source ecosystem.
- Data lakes are deployed on mature cloud storage sub-systems, allowing them to benefit from the scalability, monitoring ease of deployment, and the low storage costs associated with these systems.
- Unlike data warehouses, data lakes support all data types, including semi-structured and unstructured data, enabling workloads such as media processing.
- Because of their high throughput ingress channels, they are very well suited for streaming use cases, such as the ingestion of IoT sensor data, media streaming, or Web clickstreams.

Challenges with Data Lake

As data lakes become more popular and widely used, organizations started to recognize some challenges with traditional data lakes.

- While the underlying cloud storage is relatively inexpensive, building and maintaining an effective data lake requires expert skills.
- While it is easy to ingest data into the data lake in its raw form, transforming the data into a form that can deliver business values can be very expensive.
- Traditional data lakes have low query performance, so they cannot be used for interactive queries.
 - As a result, the organization's data teams must still transform and load the data into something like a data warehouse, resulting in an extended time to value.
 - This resulted in a data lake + warehouse architecture. This architecture continues to be dominant in the industry
- Data lakes do not offer any kind of transactional guarantees.

Challenges with Data Lake


- Data lakes have typically used a “schema on read” strategy, where data can be ingested in any format without schema enforcement. This lack of schema enforcement can result in data quality issues, allowing the pristine data lake to become a “data swamp”.
- Data files can only be appended to, leading to expensive re-writes of previously written data to make a simple update. This leads to an issue referred to as the “small file problem”.
- **Lakehouse** combines the strengths, and addresses the weaknesses of both ‘data warehouse’ and ‘data lake’ technologies.

Lakehouse


- **Lakehouse** as a system which merges the flexibility, low cost, and scale of a data lake with the data management and ACID transactions of data warehouses, addressing the limitations of both.
- Like *data lakes*, the lakehouse architecture leverages the low-cost cloud storage systems, with the inherent flexibility and horizontal scalability that comes with those systems. It also continues to leverage the data formats like Parquet and AVRO, which enable the integration with machine learning and AI systems.
- Like *data warehouses*, a data lakehouse will enable ACID transactions on data assets stored on these low-cost storage systems, enabling the reliability that used to be the exclusive domain of relational database systems.

Lakehouse

- With the data lakehouse architecture, we no longer need to have a copy of our data in the data lake, and another copy in some type of data warehouse storage.
- We can serve up our data directly from the Data Lake with comparable performance to a classical data warehouse.



Evolution



Delta Lake


- Delta Lake (*delta.io*) is an open-source storage framework that enables building a Lakehouse architecture with compute engines including Spark, PrestoDB, Flink, Trino, and Hive and APIs for Scala, Java, Rust, and Python.
- Delta Lake is a file-based metadata, caching and indexing layer on top of a data lake storage that provides an abstraction to serve ACID transactions (CRUD operations – INSERT, UPDATE, DELTE and MERGE/UPSERT).
- Delta Lake provides ACID transactions, scalable meta data handling, a unified process model that spans batch and streaming, full audit history, and support for SQL DML statements.
- It can run on existing data lakes and is fully compatible with several processing engines, including Apache Spark.

Delta Lake Features

Features of Delta lake are:


- Transactional ACID guarantees
- Full DML support
- Audit History
- Unification of batch and streaming into one processing model
- Schema enforcement and evolution
- Rich metadata support and scaling

Delta Lake Architecture



Medallion Architectural Pattern

A medallion (aka multi-hop) architecture is a data design pattern used to logically organize data in a lakehouse, with the goal of incrementally and progressively improving the structure and quality of data as it flows through each layer of the architecture (from Bronze \Rightarrow Silver \Rightarrow Gold layer tables).





Working with Delta Tables

Databricks Delta Tables

- A Delta table is a managed table that uses the Databricks Delta Optimistic Concurrency Control feature to provide full ACID transactions.
- Databricks Delta tables support all common operations, such as:
 - CRUD (create, read, update, delete)
 - Upsert (update or insert)
 - Merge (upsert a set of records)
 - Delete from the table where ... (delete based on predicate)
- In addition, Databricks Delta tables support the following:
 - Time travel – rollback data as of any point in time (based on timestamp or version)
 - Uncommitted reads - see changes that have not yet been committed
 - Optimistic concurrency control - prevents dirty reads and writes




Structured Streaming & Auto Loader

Spark Structured Streaming

- Structured Streaming is a high-level API for stream processing in Spark (version 2.2 onwards) . It is a near-real time processing engine that offers end-to-end fault tolerance with exactly-once processing guarantees using familiar Spark APIs.
- Structured Streaming lets you express computation on streaming data in the same way you express a batch computation on static data. The Structured Streaming engine performs the computation incrementally and continuously updates the result as streaming data arrives.
- The best thing about Structured Streaming is that it allows you to rapidly and quickly get value out of streaming systems with virtually no code changes. It also makes it easy to reason about because you can write your batch job as a way to prototype it and then you can convert it to a streaming job.


Spark Structured Streaming



Structured Streaming Processing Model
Users express queries using a batch API; Spark incrementalizes them to run on streams

Databricks Auto Loader

- **Autoloader** is an Optimized File Source that automatically performs incremental data loads from your Cloud storage as it arrives into the Delta Lake Tables.
- Autoloader presents a **new Structured Streaming Source** called [cloudFiles](#).
- With the DBFS paths or direct paths to the data source as the input, it automatically sets up file notifications that subscribe to file events to handle new files on arrival with an option to process the existing ones in the directory.



Databricks Auto Loader



Auto Loader file detection modes

Databricks Autoloader supports two methods to detect new files in your Cloud storage namely:

- **Directory Listing:**
 - This approach is useful for cases when only a few files are required to be streamed regularly. Here, the new files are recognized from listing the input directory.
 - With just access to your Cloud Storage data, you can swiftly enable your Databricks Autoloader Streams. From the beginning, Databricks Autoloader automatically detects if the input directory is good for Incremental Listing.
 - You have the option to explicitly choose between the Incremental Listing or Full Directory Listing by setting `cloudFiles.useIncrementalListing` as true or false.
- **File Notification**
 - Using the Cloud services like Azure Event Grid and Queue Storage services, AWS SNS and SQS or GCS Notifications, and Google Cloud Pub/Sub services, it subscribes to file events in the input directory.

Key Features of Databricks Autoloader

Using Autoloader can simplify your Data Ingestion process providing the following benefits to you:

- **Scalability:** Databricks Autoloader can track billions of files by leveraging the Cloud Services and RockDB without the need to list all the files in your directory.
- **Cost-Effective:** The notification mode for file detection eliminates the need for a directory list, thereby reducing costs. The cost of file discovery also directly depends on the number of files that are ingested rather than the number of directories in which the files arrive.
- **Ease of Use:** The Databricks Autoloader sets up the notification mode and message queue services to perform Incremental Data Load. You also don't require to track files or manage any state information on what files have arrived.
- **Schema Inference and Evolution:** For cases when there are schema drifts such as new columns, Databricks Autoloader will manage it and notify you whenever schema changes.



Lakeflow Declarative Pipelines

Lakehouse Declarative Pipelines (LDP)

- LDP (erstwhile Delta Live Tables) is a **declarative framework for building data processing pipelines** that are reliable, maintainable, and testable.
 - You define the transformations to perform on your data and LDP manages task orchestration, cluster management, monitoring, data quality, and error handling.
 - Instead of defining your data pipelines using a series of separate Apache Spark tasks, you define streaming tables and materialized views that the system should create and keep up to date.
 - Delta Live Tables manages how your data is transformed based on queries you define for each processing step. You can also enforce data quality with Delta Live Tables expectations, which allow you to define expected data quality and specify how to handle records that fail those expectations.

Why Delta Live Tables ?

- Delta Live Tables (DLT) makes it easy to build and manage reliable batch and streaming data pipelines that deliver high-quality data on the Databricks Lakehouse Platform.
- DLT helps data engineering teams simplify ETL development and management with declarative pipeline development, automatic data testing, and deep visibility for monitoring and recovery.
- Features
 - Easily build and maintain data pipelines
 - Automatic data quality testing
 - Cost-effective streaming through efficient compute auto-scaling
 - Deep visibility for pipeline monitoring and observability

Delta Live Tables datasets

Dataset type	How are records processed through defined queries?
Streaming table <small>STREAMING LIVE TABLE</small>	Each record is processed exactly once. This assumes an append-only source.
Materialized view <small>LIVE TABLE</small>	Records are processed as required to return accurate results for the current data state. Materialized views should be used for data sources with updates, deletions, or aggregations, and for change data capture processing (CDC).
View	Records are processed each time the view is queried. Use views for intermediate transformations and data quality checks that should not be published to public datasets.

DLT datasets – Streaming Table

- A streaming table is a Delta table with extra support for streaming or incremental data processing. They allow you to process a growing dataset, handling each row only once.
 - Because most datasets grow continuously over time, streaming tables are good for most ingestion workloads.
 - Streaming tables are optimal for pipelines that require data freshness and low latency.
 - Streaming tables can also be useful for massive scale transformations, as results can be incrementally calculated as new data arrives, keeping results up to date without needing to fully recompute all source data with each update.
- Streaming tables are designed for data sources that are append-only.

DLT datasets – Materialized View

- A materialized view (or live table) is a view where the results have been pre-computed.
- Materialized views are refreshed according to the update schedule of the pipeline in which they're contained.
- Materialized views can handle any changes in the input. Each time the pipeline updates, query results are recalculated to reflect changes in upstream datasets that might have occurred because of compliance, corrections, aggregations, or general CDC.
- Delta Live Tables implements materialized views as Delta tables, but abstracts away complexities associated with efficient application of updates, allowing users to focus on writing queries.

DLT datasets – View

- All views in Databricks compute results from source datasets as they are queried, leveraging caching optimizations when available.
- Delta Live Tables does not publish views to the catalog, so views can be referenced only within the pipeline in which they are defined.
- Views are useful as intermediate queries that should not be exposed to end users or systems.
- Databricks recommends using views to enforce data quality constraints or transform and enrich datasets that drive multiple downstream queries.

DLT Pipeline


- A **pipeline** is the main unit used to configure and run data processing workflows with Delta Live Tables.
- A pipeline contains materialized views and streaming tables declared in Python or SQL source files.
- Delta Live Tables infers the dependencies between these tables, ensuring updates occur in the correct order. For each dataset, Delta Live Tables compares the current state with the desired state and proceeds to create or update datasets using efficient processing methods.
- The settings of Delta Live Tables pipelines fall into two broad categories:
 - Configurations that define a collection of notebooks or files (known as source code or libraries) that use Delta Live Tables syntax to declare datasets.
 - Configurations that control pipeline infrastructure, how updates are processed, and how tables are saved in the workspace.

Monitor and enforce data quality

- You can use ***expectations*** to specify data quality controls on the contents of a dataset.
 - Unlike a CHECK constraint in a traditional database which prevents adding any records that fail the constraint, expectations provide flexibility when processing data that fails data quality requirements.
 - This flexibility allows you to process and store data that you expect to be messy and data that must meet strict quality requirements.

Some important points to note

- DLT notebooks can not be run interactively. They are intended to be attached to a DLT pipeline. You can attach more than one notebook to a pipeline.
- Only Python and SQL are supported. Each notebook can have only one language – either python or SQL. We can not mix both
- We can not use most magic commands like %run
- Requires a premium Databricks workspace
- Automatically generated two clusters – execution and maintenance clusters.




Unity Catalog

Unity Catalog

Unity Catalog is a unified governance solution for data and AI assets on the Databricks Lakehouse.

Unity Catalog provides centralized access control, auditing, lineage, and data discovery capabilities across Databricks workspaces.



- **Unity Catalog is supported on clusters that run Databricks Runtime 11.3 LTS or above.**

Unity Catalog features


- **Define once, secure everywhere:** Unity Catalog offers a single place to administer data access policies that apply across all workspaces.
- **Standards-compliant security model:** Unity Catalog's security model is based on standard ANSI SQL and allows administrators to grant permissions in their existing data lake using familiar syntax, at the level of catalogs, databases (also called schemas), tables, and views.
- **Built-in auditing and lineage:** Unity Catalog automatically captures user-level audit logs that record access to your data. Unity Catalog also captures lineage data that tracks how data assets are created and used across all languages.
- **Data discovery:** Unity Catalog lets you tag and document data assets, and provides a search interface to help data consumers find data.

Unity Catalog object model

In Unity Catalog, the hierarchy of primary data objects flows from metastore to table or volume:

- **Metastore:** The top-level container for metadata. Each metastore exposes a three-level namespace (*catalog.schema.table*) that organizes your data.
- **Catalog:** The first layer of the object hierarchy, used to organize your data assets.
- **Schema:** Also known as databases, schemas are the second layer of the object hierarchy and contain tables and views.
- **Volume:** Volumes sit alongside tables and views at the lowest level of the object hierarchy and provide governance for non-tabular data.
- **Table:** At the lowest level in the object hierarchy are tables and views.

Unity Catalog object model



Metastore

- A metastore is the top-level container of objects in Unity Catalog. It stores metadata about data assets (tables and views) and the permissions that govern access to them.
- Databricks account admins should create one metastore per region in which they operate and assign them to Databricks workspaces in the same region.
- For a workspace to use Unity Catalog, it must have a Unity Catalog metastore attached.
- Each metastore is configured with a managed storage location in an S3 bucket in your AWS account.

Managed Storage

- When an account admin creates a metastore, they must associate a storage location in an S3 bucket in your AWS account to use as managed storage.
- Unity Catalog also allows users to associate managed storage locations with catalogs and schemas.
- Managed storage has the following properties:
 - Managed tables and managed volumes store data and metadata files in managed storage.
 - Managed storage cannot overlap with external tables, external volumes, or other managed storage.

Catalog

- A catalog is the first layer of Unity Catalog's three-level namespace. It's used to organize your data assets.
- Users can see all catalogs on which they have been assigned the **USE CATALOG** data permission.
- All users have the **USE CATALOG** permission on the **main** catalog. The main catalog is intended for organizations that are just getting started with Unity Catalog. As you add users and data, you should add catalogs to maintain a data hierarchy that enables efficient control over access.

Schema

- A schema (also called a database) is the second layer of Unity Catalog's three-level namespace.
- A schema organizes tables and views. Users can see all schemas on which they have been assigned the **USE SCHEMA** permission, along with the **USE CATALOG** permission on the schema's parent catalog.
- To access or list a table or view in a schema, users must also have **SELECT** permission on the table or view.

Tables

- A table resides in the third layer of Unity Catalog's three-level namespace.
 - To create a table, users must have CREATE and USE SCHEMA permissions on the schema, and they must have the USE CATALOG permission on its parent catalog.
 - To query a table, users must have the SELECT permission on the table, the USE SCHEMA permission on its parent schema, and the USE CATALOG permission on its parent catalog.
- A table can be managed or external.
 - **Managed Table** : Managed tables are the default way to create tables in Unity Catalog. Unity Catalog manages the lifecycle and file layout for these tables. You should not use tools outside of Databricks to manipulate files in these tables directly. Managed tables **always use the Delta table format**.
 - **External Table**: External tables are tables whose data lifecycle and file layout are not managed by Unity Catalog. Use external tables to register large amounts of existing data in Unity Catalog, or if you require direct access to the data using tools outside of Databricks clusters or Databricks SQL warehouses. When you drop an external table, Unity Catalog does not delete the underlying data. External tables can use the following file formats: DELTA, CSV, JSON, AVRO, PARQUET, ORC, TEXT.

How to set up Unity Catalog for my organization?

To set up Unity Catalog for your organization, you do the following:

- Configure an S3 bucket and IAM role that Unity Catalog can use to store and access data in your AWS account.
- Create a metastore for each region in which your organization operates.
- Attach workspaces to the metastore.
 - Each workspace will have the same view of the data you manage in Unity Catalog.
- If you have a new account, add users, groups, and service principals to your Databricks account. Create and grant access to catalogs, schemas, and tables.

Create a cluster or SQL warehouse

Before you can start creating tables and assigning permissions, you need to create a compute resource to run your table-creation and permission-assignment workloads.

Tables defined in Unity Catalog are protected by fine-grained access controls. To ensure that access controls are enforced, Unity Catalog requires compute resources to conform to a secure configuration. Non-conforming compute resources cannot access tables in Unity Catalog.

Databricks provides two kinds of compute resources:

- **Clusters**, which are used for executing commands in Databricks notebooks and running jobs.
- **SQL warehouses**, which are used for executing queries in SQL Editor.

You can use either of these compute resources to work with Unity Catalog.



Delta Sharing

Delta Sharing

- Delta Sharing in Databricks is a secure data sharing platform that lets you share data and AI assets in Databricks with users outside your organization, regardless of whether they use Databricks.
- Delta Sharing is the basis for *Databricks Marketplace*, an open forum for exchanging data products, and *Clean Rooms*, a secure and privacy-protecting environment where multiple parties can work together on sensitive enterprise data.
- Delta Sharing is available as fully managed service inside Databricks platform for workspaces enabled with Unity Catalog.

Delta Sharing Protocols


- **Databricks-to-Databricks sharing protocol**
 - lets you share data and AI assets from your Unity Catalog-enabled workspace with users who also have access to a Unity Catalog-enabled Databricks workspace.
- **Databricks open sharing protocol**
 - lets you share tabular data that you manage in a Unity Catalog-enabled Databricks workspace with users on any computing platform.
- **Customer-managed open-source Delta Sharing server**
 - lets you share from any platform to any platform, whether Databricks or not.

Delta Sharing – Key concepts

- **Share**
 - is a read-only collection of tables and table partitions that a provider wants to share with one or more recipients. If your recipient uses a Unity Catalog-enabled Databricks workspace, you can also include notebook files, views (including dynamic views that restrict access at the row and column level), Unity Catalog volumes, and Unity Catalog models in a share.
- **Provider**
 - is an entity that shares data with a recipient.
 - If you are a provider and you want to take advantage of the built-in Databricks Delta Sharing server and manage shares and recipients using Unity Catalog, you need at least one Databricks workspace that is enabled for Unity Catalog.
- **Recipient**
 - is an entity that receives shares from a provider.

Databricks-to-Databricks Delta Sharing model

- A data *recipient* gives a data *provider* the unique *sharing identifier* for the Databricks Unity Catalog metastore that is attached to the Databricks workspace.
- The data provider creates a *share* in the provider's Unity Catalog metastore.
 - This named object contains a collection of tables, views, volumes, and notebooks registered in the metastore.
- The data provider creates a recipient object in the provider's UC metastore.
- The data provider grants the recipient access to the share.
- The share becomes available in the recipient's Databricks workspace.



Lakehouse Federation

Lakehouse Federation

- **Lakehouse Federation** is the query federation platform for Databricks.
 - The term query **federation** describes a collection of features that enable users and systems to run queries against multiple data sources without needing to migrate all data to a unified system.



Lakehouse Federation

There are two types of federation:

- Query federation
- Catalog federation



Query vs. Catalog Federation

Query Federation	Catalog Federation
Direct JDBC connection to external databases (eg. MySQL, SQL Server, Redshift)	Connects to external catalogs (AWS Glue, Hive Metastore, Snowflake)
Query executed on the external database	Query executed on Databricks compute
Performance/cost depend on source system	More cost-effective and performance-optimized
Can join with Lakehouse data	Can join with Lakehouse data
Best for ad-hoc or PoC	Best for migration or hybrid models

Query vs. Catalog Federation

Query Federation	Catalog Federation
Unity Catalog enabled	Unity Catalog enabled
Create connection (JDBC + credentials)	Create connection to external catalog
Create foreign catalog	Create storage credential + external location
Grant privileges	Create foreign catalog
Run queries → pushed down to external DB	Grant privileges
	Run queries → executed on object storage



Databricks Asset Bundles


Databricks Asset Bundles (DAB)

- Databricks Asset Bundles are a tool to facilitate the adoption of software engineering best practices, including source control, code review, testing and CI/CD, for your data and AI projects.
- Bundles provide a way to include metadata alongside your project's source files and make it possible to describe Databricks resources such as jobs and pipelines as source files (yml files).
- A bundle is an end-to-end definition of a project, including how the project should be structured, tested, and deployed.

Databricks Asset Bundles (DAB)

- Your bundle project's collection of source files and metadata is deployed as a single bundle to your target environment.
- A bundle includes the following parts:
 - Required cloud infrastructure and workspace configurations
 - Source files, such as notebooks and Python files, that include the business logic
 - Definitions and settings for Databricks resources, such as Lakeflow Jobs, Lakeflow Declarative Pipelines, Dashboards etc.
 - Unit tests and integration tests

Databricks Asset Bundles (DAB)



How do Bundles work?

- Bundle metadata is defined using YAML files that specify the artifacts, resources, and configuration of a Databricks project. The Databricks CLI can then be used to validate, deploy, and run bundles using these bundle YAML files. You can run bundle projects from IDEs, terminals, or within Databricks directly.
- Bundles can be created manually or based on a template. The Databricks CLI provides default templates for simple use cases, but for more specific or complex jobs, you can create custom bundle templates to implement your team's best practices and keep common configurations consistent.

What do I need to install?

- Databricks Asset Bundles are a feature of the Databricks CLI. You build bundles locally, then use the Databricks CLI to deploy your bundles to target remote Databricks workspaces and run bundle workflows in those workspaces from the command line.
- Databricks CLI installation instructions:
 - <https://docs.databricks.com/aws/en/dev-tools/cli/install>



Databricks Local Development

Local development environment setup

The following are the step for setting up local development environment for Databricks on Visual Studio Code IDE.

- Install Python
- Install Visual Studio Code
- Install Databricks Plugin in VSCode
- Create Databricks Project
- Configure Access to Databricks
- Run code on Databricks cluster

Which python version to use?


Databricks Connect version	Compute type	Compatible Python version
16.0.0 to 16.4.0	Serverless	There is no compatible Python version. Upgrade to Databricks Connect 16.4.1 or above.
16.0.x to 16.3.x	Cluster	3.12
15.1.0 to 15.4.9	Serverless	There is no compatible Python version. Upgrade to Databricks Connect 15.4.10 or above.
15.1.x to 15.3.x	Cluster	3.11
14.0.x to 14.2.x	Cluster	3.10
13.0.x to 13.2.x	Cluster	3.10

Databricks Connect

- Databricks Connect is a client library for the Databricks Runtime that allows you to connect to Databricks compute from IDEs such as Visual Studio Code, PyCharm, and IntelliJ IDEA, notebooks and any custom application, to enable new interactive user experiences based on your Databricks Lakehouse.
- Databricks Connect is available for the following languages:
 - Databricks Connect for Python
 - Databricks Connect for R
 - Databricks Connect for Scala

Why Databricks Connect?

- Using Databricks Connect, you can write code using Spark APIs and run them remotely on Databricks compute instead of in the local Spark session.
 - **Interactively develop and debug from any IDE.**
 - **Build interactive data apps.**



Thank You